Physics-Informed Machine-Learning Inversion of Chandrayaan-3 RAMBHA-LP Langmuir Probe Sweeps

Sunil Kumar (M.Tech CSE, IIIT Hyderabad) · Supervisor: Dr. Rama Chandra Pillutla

The RAMBHA-LP Langmuir probe on the Chandrayaan-3 Vikram lander recorded current–voltage sweeps of the near-surface lunar plasma near the south pole between 24 August and 3 September 2023. Electron density, electron temperature and plasma potential are conventionally extracted by piecewise curve fitting, which fails on 44% of the 29,862 sweeps in the public rawB archive. We trained a physics-informed neural network on 40,000 synthetic sweeps from an orbital-motion-limited forward model, with noise and parameter ranges measured from the real archive, and a loss that also penalises failure to reconstruct the input curve. 28% of the sweeps are clipped by the instrument and the network's values on them are artefacts, so they are excluded; on the rest, the network recovers 5,410 of the 5,936 unclipped classical failures, raising usable coverage from 55.6% to 73.7%. With no ground truth on the Moon, we validate against an independent classical fit (no systematic density offset; 4.7% per sweep over 16,596 sweeps), against the classical value on the other half of the same sweep for 2,410 recovered sweeps (5.2%), and by self-consistency between the rising and falling halves of each commanded sweep (2.4%, r = 0.993, 10,365 pairs); the daily medians also follow the mission-long density rise from 246 to 500 cm^{-3}, as the classical fit does. Median values agree with the published ranges for this instrument, except for density on the first three days. Physics-informed training worsens accuracy on synthetic data, yet on real sweeps it makes the two halves of each sweep 29% more consistent and removes a 6% density offset against the classical fit, showing that simulation alone cannot judge sim-to-real choices. A two-population extension with a working negative control finds no evidence for a second electron population, in agreement with the instrument team, and shows that the standard averaged-sweep diagnostic for multiple populations is not valid evidence on its own.