Reading the Moon's Plasma with a Neural Network

A physics-informed inversion of Chandrayaan-3 RAMBHA-LP Langmuir probe sweeps

Sunil Kumar
M.Tech Computer Science, IIIT Hyderabad
Independent Study · Supervisor: Dr. Rama Chandra Pillutla

Abstract

The RAMBHA-LP Langmuir probe on the Chandrayaan-3 Vikram lander recorded current–voltage sweeps of the near-surface plasma near the lunar south pole between 24 August and 3 September 2023. Electron density, electron temperature and plasma potential are usually obtained from such sweeps by piecewise line fitting, which fails on 44% of the 29,862 sweeps in the public archive. More than half of those failures, and 30% of all sweeps, were clipped by the instrument, and neither method can read them. This study trains a physics-informed neural network on 40,000 synthetic sweeps from an orbital-motion-limited probe model, using a loss that also requires the predicted parameters to rebuild the input curve. The network recovers 5,115 of the 5,641 unclipped sweeps that the classical fit rejects (91%), raising the usable part of the archive from 50.7% (the classical fit on unclipped sweeps) to 67.9%. Since no ground truth exists on the Moon, the results are checked against the classical fit (no meaningful density offset; 4.4% typical difference per sweep), against the classical value on the other half of the same sweep for 2,157 recovered sweeps (5.3%), by comparing the rising and falling halves of each sweep (2.4%), and against values published by the instrument team. Comparing the six channel and resistor setups, which sample the same plasma seconds apart, shows that two of them read density low, by about 17% and, before 27 August, by about half; both methods share these offsets, and 51% of the archive is usable without any setup correction. The physics term costs accuracy on simulated data but makes real densities more repeatable and removes a density offset against the classical fit. A two-population extension, tested against a single-population control, finds no evidence for a second electron population, and the averaged-sweep diagnostic used to infer such populations is shown to be unreliable when sweep temperatures vary. Electron temperature, about 20% on simulated data and 25% per sweep against the classical fit, is limited by the forward model near the plasma potential rather than by the network.

1. Introduction

The Moon has no global magnetic field and only a tenuous exosphere, so its sunlit surface is exposed directly to solar ultraviolet radiation and the solar wind. Photoemission from the regolith charges the dayside surface positively and fills the region just above it with a photoelectron sheath, whose structure controls surface charging and the levitation of fine dust (Poppe & Horányi 2010). The electron density, electron temperature and plasma potential of this near-surface layer therefore matter both for understanding the surface–plasma interaction and for the electrical environment in which landers and surface instruments operate. Most earlier measurements of the lunar ionosphere were made from orbit by radio occultation (Manju et al. 2026), and measurements from the surface itself have been rare.

Chandrayaan-3 provided a new opportunity. Its Vikram lander touched down near the lunar south pole on 23 August 2023 and carried RAMBHA-LP (Radio Anatomy of Moon Bound Hypersensitive ionosphere and Atmosphere - Langmuir Probe), a spherical probe mounted on a boom above the lander deck. From these measurements Manju et al. (2025) reported electron densities of 380–600 cm^{-3} and electron temperatures of 3,000–8,000 K near the landing site, the first in situ near-surface plasma measurements from the south polar region. A second study used the lander's hop experiment on 2 September 2023 and reported a fall in mean density from about 478 cm^{-3} before the hop to about 26 cm^{-3} during and immediately after it (Manju et al. 2026). The sweep data are publicly released through the ISRO Science Data Archive (ISRO PRADAN), which makes an independent re-analysis of the full archive possible.

A Langmuir probe infers plasma parameters from the current it collects while its bias voltage is swept. The theory goes back to Mott-Smith & Langmuir (1926), whose orbital-motion-limited (OML) description remains the basis for analysing small spherical and cylindrical probes (Hutchinson 2002). Below the plasma potential the electron current rises exponentially with a slope set by the electron temperature; above it the current grows more slowly as the collecting sheath expands, and its level scales with the electron density. Probes of this kind have flown widely. The Langmuir Probe and Waves instrument on MAVEN measures electron density and temperature in the Martian ionosphere (Ergun et al. 2015), and probe sweeps on Rosetta were used to quantify photoelectron emission and track the solar ultraviolet flux at comet 67P (Johansson et al. 2017).

The usual analysis fits straight lines to the logarithm of the current in the electron retardation region and to the saturation region, and reads the three parameters from those fits (Hutchinson 2002). It works well on clean, well-sampled sweeps, but the result depends on where the fitting windows are placed, and it becomes unreliable when the curve is noisy, poorly sampled in the exponential part, or modified by photoemission from the probe and its support (Johansson et al. 2017). On RAMBHA-LP, the classical fit fails on 44% of the sweeps (Section 1.1), so a large part of a dataset that cannot be repeated is at present unused.

Machine-learning regression offers another route. Instead of fitting selected parts of the curve, a model trained on many examples maps the whole sweep to plasma parameters. Liu et al. (2023) trained radial basis function and neural network regressors on synthetic probe currents from three-dimensional kinetic simulations and applied them to multi-needle Langmuir probe data from the NorSat-1 satellite, showing that simulation-trained inference can be used on flight data. For the Moon the main difficulty is that no independent measurement of the near-surface plasma exists, so training labels can only come from a physical model. Physics-informed neural networks address part of this problem by adding the governing equations to the training loss, so that predictions must also be consistent with the physics (Raissi et al. 2019). Monte-Carlo dropout can attach an uncertainty to each prediction at little extra cost (Gal & Ghahramani 2016), which is valuable when results cannot be checked against a known answer.

This study applies these ideas to the RAMBHA-LP archive. A physics-informed neural network is trained only on synthetic sweeps from an OML forward model and is then used to invert the real sweeps, including those that the classical fit rejects. Because there is no ground truth, as much weight is given to validation as to the inversion itself: agreement with the classical method, consistency between the rising and falling halves of each sweep, a check against the mission timeline, and comparison with published values. The same framework is then used to test whether the near-surface plasma contains a second, hotter electron population, and to find out what limits the accuracy of the electron temperature on this instrument.

1.1 Why the standard method struggles

Applied to this archive, the classical fit succeeds on 16,596 of 29,862 sweeps (55.6%). Table 1 lists the main reasons for failure; a further 136 sweeps fail for smaller reasons. It also shows how many of each kind were clipped by the instrument (Section 2.5): 7,625 of the 13,266 failures, including all but one of the flat-top failures. Clipped sweeps cannot be read by either method, so the failures this study can address are the other 5,641, 19% of the archive. The failure count depends on the width of the fit window (Section 2.4): of the unclipped sweeps, 16,219 pass with a 0.8 V window and 13,023 with 1.5 V, against 15,147 with 1.0 V, while the median density of the sweeps that pass at all three widths changes by at most about 0.5%.

Table 1. *Reasons the classical fit failed, with the number of sweeps affected and how many of them the instrument clipped.*

Why the fit failed	Sweeps	Clipped	What it means
Flat top has the wrong slope	2,369	2,368	the instrument clipped the top of the curve
Too few points in the steep region	2,347	113	the exponential is too narrow to fit
Fit quality below threshold (R^2 under 0.90)	4,945	2,639	curve shape does not match the model; about half are clipped
Plasma potential outside the sweep range	3,469	2,482	the saturation fit was distorted, mostly by clipping

Recovering the 5,641 unclipped sweeps that the classical fit discards, from a dataset that cannot be repeated, is the problem this study addresses.

1.2 Objectives

The objectives of this study are:

- To recover electron density, electron temperature and plasma potential from RAMBHA-LP sweeps that the classical fit cannot process, and so increase the usable fraction of the archive.
- To validate the recovered parameters without ground truth, using an independent method, the instrument's own repeated measurements and published results.
- To measure, rather than assume, how the physics-informed loss term affects performance on simulated and on real data.
- To test, with appropriate controls, whether the archive requires a second electron population.
- To identify the factor that limits electron temperature accuracy for this instrument.

2. Data and methodology

2.1 Instrument and data

The data were obtained from the ISRO Science Data Archive, PRADAN, operated by the Indian Space Science Data Centre (ISRO PRADAN). RAMBHA-LP products are distributed in the Planetary Data System version 4 (PDS4) format at three processing levels, labelled raw, rawA and rawB. The raw level holds only the bias voltage and the probe output voltage, which must be converted to current using the instrument's resistor and gain tables. In rawA (Level-0A) this conversion has been done and the current is given in amperes. In rawB (Level-0B) a photoemission current, estimated with the method of Johansson et al. (2017), has also been subtracted. This study uses rawB because it is the archive's calibrated, photoemission-corrected product and needs no further gain conversion. Within each file the difference between rawA and rawB is a constant offset, which the baseline subtraction described in Section 2.2 removes, so the choice of level does not change the curve shape that the network sees.

The release used is LTA_RAMBHA-LP_August2024_final, and its rawB level contains 144 science (sweep) files with their ops files. The science files hold the probe current, and the ops files record the commanded sweep sequence, including the number of samples in each sweep. Together they cover 24 August to 3 September 2023. In sweep mode the probe bias was stepped from −12 V to +12 V in 0.1 V steps and back again. The current was recorded on two ADC channels, with three probe resistor settings available in the front end; the six combinations are used in turn, one per sweep segment.

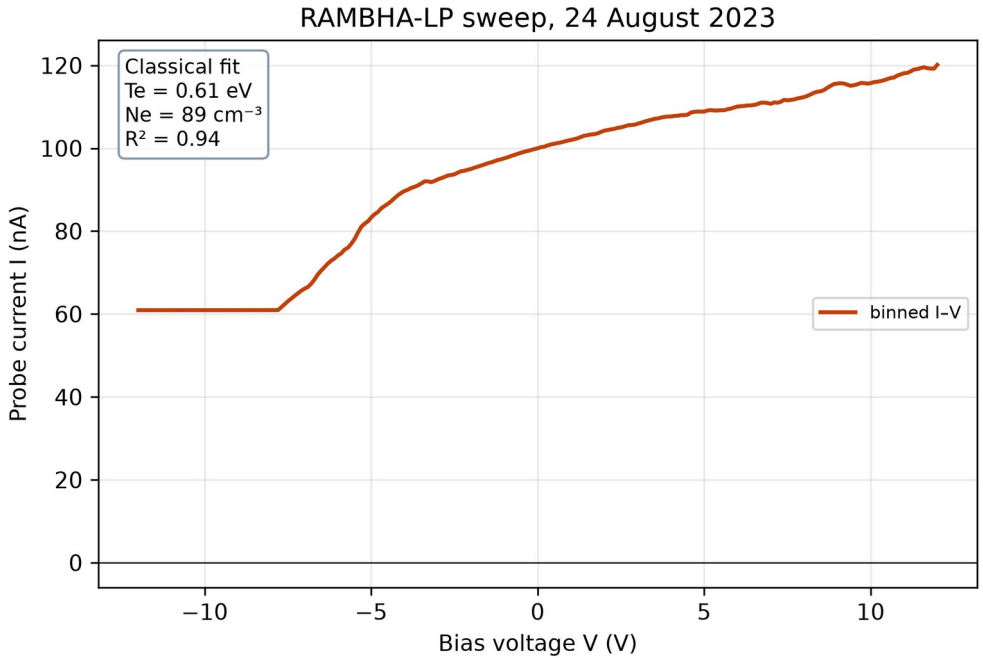


Figure 1. *One real RAMBHA-LP sweep from 24 August 2023: binned probe current against bias voltage. The flat floor on the left, the steep rise and the slower climb on the right are the parts of the curve that the three plasma parameters are read from.*

2.2 Preprocessing

Each science file was paired with its ops file, and the ops sample counts were used to cut the continuous science record into individual commanded sweeps. A commanded sweep is a triangle in which the bias rises from −12 V to +12 V and then returns. Each triangle was split into its rising and falling halves, which were treated as separate sweeps. This turned 14,931 commanded sweeps into 29,862 sweeps and also provides the paired measurements used for the self-consistency test in Section 3.2.2.

Within each half, the samples recorded at each voltage step were averaged, which reduces random noise without changing the shape of the curve. The averaged curves were resampled onto a common grid of 240 voltage points, so that every sweep has the same length as the network input. All sweeps were then brought to a common orientation and the ion-current baseline, estimated as the median of the lowest 10% of the curve, was subtracted; the classical fit uses its own estimate (Section 2.4). Finally, the currents were expressed on one fixed current scale for the whole archive rather than normalised sweep by sweep, so that the absolute current level, which carries the density information, is preserved.

2.3 Physics model and synthetic training data

No independent measurement of the near-surface plasma exists, so the network cannot be trained on measured labels. Labelled examples were instead generated with a physics model that predicts the current curve for a given plasma. The network therefore learns to invert the model rather than the plasma itself, and any error in the model passes into its results. This limitation recurs in Section 3.3.

The forward model follows orbital-motion-limited probe theory for a sphere (Mott-Smith & Langmuir 1926). The electron current at the plasma potential is

$$I_{e0} = e\,N_e\,A\sqrt{\frac{e\,T_e}{2\pi\,m_e}},\quad A = 4\pi r^2 \qquad (1)$$

where e is the electron charge, Ne the electron density, Te the electron temperature in eV, me the electron mass and r = 2.5 cm the probe radius. Below the plasma potential Vp the current falls off exponentially, and above it the current grows with a sheath-expansion exponent α, set to 0.40 from fits to real sweeps:

$$I_e(V) = I_{e0}\exp\left(\frac{V - V_p}{T_e}\right),\quad V < V_p \qquad (2)$$

$$I_e(V) = I_{e0}\left(1 + \frac{V - V_p}{T_e}\right)^{\alpha},\quad V \geq V_p \qquad (3)$$

The exponent was chosen from the data alone. Across the 4,491 unclipped sweeps that the classical fit (Section 2.4) passes at every value tried from 0.2 to 1.0, the saturation branch is described best at α = 0.40 (median R^2 0.983, against 0.956 at 0.5 and 0.884 at the textbook value of 1). It was first selected during development partly because it reproduced the published values for 2 September, which is why this data-only check matters for Section 3.2.5. The best value is 0.40 on the 49.9 kΩ and 1 MΩ resistor settings but 0.50 on the 20 MΩ setting, where all the unclipped sweeps come from channel 1, so a dependence on the measurement chain cannot be excluded; that setup also reads density low (Section 3.2.6).

The synthetic data consist of 40,000 training sweeps and a separate test set of 4,000 sweeps. For each sweep, the electron density, electron temperature and plasma potential were drawn at random from 2–1,000 cm^{-3}, 0.05–2 eV and −12 to +6 V respectively, and Equations 1–3 were used to compute the current on the same 240-point voltage grid as the real data. The temperature was drawn uniformly across its range; as Section 3.3.3 shows, this places a large share of the examples at temperatures above those observed at the landing site. To make the synthetic curves resemble the real ones, random noise and a linear baseline tilt were added to each curve. The noise was drawn for one of the three probe-resistor settings, log-uniformly within 0.15–0.55%, 0.15–0.45% and 0.8–4.5% of the curve's height, around the medians measured on the real sweeps (0.21%, 0.24% and 1.8%); the tilt was 0.2–10% of the height with a random sign. Clipping by the instrument (Section 2.5) was not simulated.

2.4 Classical fit

The classical fit follows the standard procedure for a spherical probe (Hutchinson 2002). The floating potential Vf is taken where the smoothed current last crosses zero, and the ion current is removed by subtracting the median current more than 2 V below Vf. The temperature is the inverse slope of ln(I) in the retardation window, which starts at Vf and is 1.0 V wide (ten samples at the 0.1 V step), ending earlier if the plasma potential is reached. Above that window the current is raised to the power 1/α, which turns Equation 3 into a straight line; its zero crossing and slope give Vp and Ie0, and Equation 1 then gives the density. A sweep fails if the temperature lies outside 0.03–5 eV, the density outside 1–100,000 cm^{-3}, the retardation fit has R^2 below 0.90, or any step cannot be completed.

The acceptance test does not require Vp to lie above Vf. In 2,331 of the 15,147 unclipped sweeps the classical fit passes (15.4%), the fitted Vp is below Vf, which is not physical. Section 3.2.1 shows that removing them changes the comparison with the network very little.

2.5 Physics-informed network

The network is trained on two requirements: its three outputs must match the known values, and feeding them back through the physics model must reproduce the input curve. This follows the physics-informed approach of Raissi et al. (2019).

The training loss combines a supervised term and a physics term:

$$L = w_s L_{\mathrm{sup}} + w_p L_{\mathrm{phys}} \qquad (4)$$

where the first term scores the predicted parameters against the known labels, the second scores how well the parameters rebuild the input curve through Equations 2 and 3, and the weights are 1 and 0.5. Each supervised error is divided by half the span of its training range, and the physics error by the mean square of the training curves, so that the weights mean what they say.

The network is a fully connected multilayer perceptron with an input of 240 current values, three hidden layers of 256, 256 and 128 units with 10% dropout, and an output layer of three values, giving 160,771 trainable parameters. The outputs are $\log_{10}$ Ne, Te and Vp; density is predicted on a logarithmic scale because it spans almost three orders of magnitude in the training set. The network was trained for 250 epochs with the loss of Equation 4, using the Adam optimiser (learning rate 10^{-3}, weight decay 10^{-5}, batch size 512) with a cosine learning-rate schedule. Dropout is kept active at prediction time, and each sweep is passed through the network 50 times; the mean of these outputs is taken as the estimate and their standard deviation as its uncertainty (Gal & Ghahramani 2016). A

real sweep that the classical method could not fit is counted as recovered only if five conditions hold: the predicted density lies between 1 and 100,000 cm^{-3}, the predicted temperature lies between 0.03 and 5 eV, the Monte-Carlo dropout spread of $\log_{10}$ Ne is below 0.5 dex, the curve rebuilt from the predicted parameters reproduces the baseline-subtracted sweep with R^2 of at least 0.90, and the sweep is not clipped. The R^2 threshold is the classical one, but it is applied to an easier quantity, the linear current over the whole sweep rather than ln(I) in a 1 V window: it accepts 99.9% of the sweeps the classical fit passes. In practice the R^2 and clipping conditions decide the outcome: of the 13,266 classical failures, 13,266 pass the density condition, 13,265 the uncertainty condition and 13,260 the temperature condition; 8,345 pass the R^2 test, and 5,115 of those are not clipped.

The instrument's front end saturates on some settings, and the top of the curve then becomes a flat plateau. A sweep is treated as clipped if its last five voltage steps, the top 0.5 V, all lie within 0.1% of its current range below its maximum; on an unsaturated curve, noise makes such a run essentially impossible. Both thresholds were checked: the rule flags about 0.1% of sweeps on the settings that never clip (49.9 kΩ, and 1 MΩ on channel 1), and at least 99.8% of real sweeps clipped on purpose (Section 3.2.3); a plateau shorter than 0.5 V is not flagged. By ADC channel and resistor setting, the clipped share is 73% on channel 2 at 1 MΩ, 100% on channel 2 at 20 MΩ, 9% on channel 1 at 20 MΩ and about 0.1% elsewhere: 9,074 sweeps (30.4%) in all, 7,625 of the classical failures and 1,449 of the classical passes. Neither method can read these curves. The physics model has no saturation; its saturation branch goes flat only through a numerical limit that is reached inside the sweep at very low temperature, and the network uses that limit to imitate the plateau. The classical fit is no better, and its standard acceptance test does not detect clipping: 1,449 clipped sweeps pass it, and on 78% of them at least half of the saturation branch it fits is the plateau. Clipped sweeps are therefore left out of both methods' usable counts (Section 3.2.3).

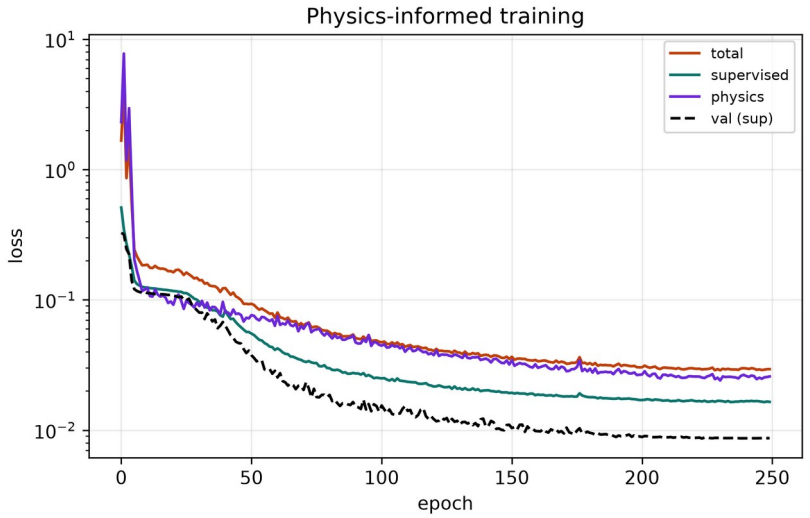


Figure 2. *Training the network. The supervised loss scores the predicted numbers against the known answer, the physics loss scores how well those numbers rebuild the curve, and the dashed line is the supervised loss on held-out simulated sweeps.*

2.6 Validation strategy

Since there is no ground truth, every check in this report is one of the three kinds listed in Table 2. The working practices behind these checks are described in Supplementary Appendix B.

Table 2. *The three kinds of check used to validate results without ground truth.*

Kind of check	What it does
Compare to an independent method	Run the classical fit on the same sweeps. It shares no code with the network, so agreement is meaningful.
Compare repeated measurements	Each commanded sweep goes up and then back down, giving two independent looks at the same plasma seconds apart. Their disagreement gives an error estimate independent of the model.
Build a control	Generate what the measurement would look like if the effect under test were absent, and compare.

3. Results and discussion

3.1 Coverage and accuracy

3.1.1 Usable fraction of the archive

Table 3. *Usable sweeps from the classical method and the network.*

	Sweeps	Share
Total sweeps	29,862	100%
Classical method succeeds	16,596	55.6%
...of which clipped (not usable)	1,449	8.7% of them
Classical fails	13,266	44.4%
...clipped by the instrument	7,625	57.5% of them
...recovered by the network	5,115	38.6% of them
Total usable after both methods	**20,262**	**67.9%**

The network does not replace the classical method. Where the classical fit works, the two show no meaningful offset (Section 3.2.1); the network extends coverage into sweeps the fitting approach cannot handle. Of the 5,641 classical failures that are not clipped, it recovers 5,115 (91%). Counting only the 20,788 unclipped sweeps, the usable share rises from 73% to 97%.

A sweep is counted as usable if it is not clipped and either the classical fit passes or the network recovers it. The clipping rule is strict: 442 clipped sweeps on channel 1 at 20 MΩ go flat only above about +11.3 V (against +4.3 V for the clipped sweeps on channel 2 at 1 MΩ) and give the same density as the unclipped sweeps on that setup; admitting them would raise the usable share to 69.3%. They are left out so that one clipping rule applies throughout. An earlier version of this analysis also counted clipped sweeps and reported 83.5%; that figure is withdrawn (Supplementary Table S1). Some usable sweeps carry a calibration offset of their setup (Section 3.2.6).

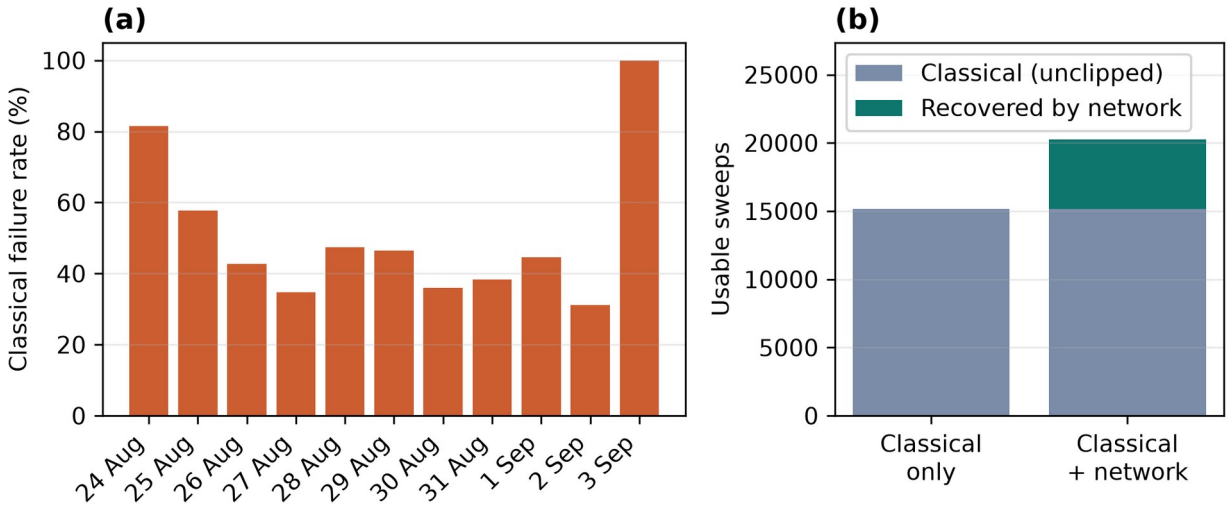


Figure 3. *(a) Share of sweeps on which the classical fitting method fails, by day; overall it fails on 44% of the 29,862 sweeps, and on every sweep of 3 September. (b) Usable sweeps with the classical method alone and with the network added, both without clipped sweeps: the network recovers 5,115 of the classical failures, taking the usable total from 15,147 to 20,262.*

3.1.2 Accuracy on simulated sweeps

Table 4 gives the error on 4,000 simulated sweeps that were not used in training.

Table 4. *Network error on 4,000 held-out simulated sweeps.*

Quantity	Typical error	Note
Electron density	0.021 dex	about 5%
Electron temperature	0.078 eV	about 20% at our typical 0.4 eV
Plasma potential	0.141 V	about a seventh of a volt

These values are a best case: they show that the network can invert the physics model, not that it can invert real data, which Section 3.2 tests.

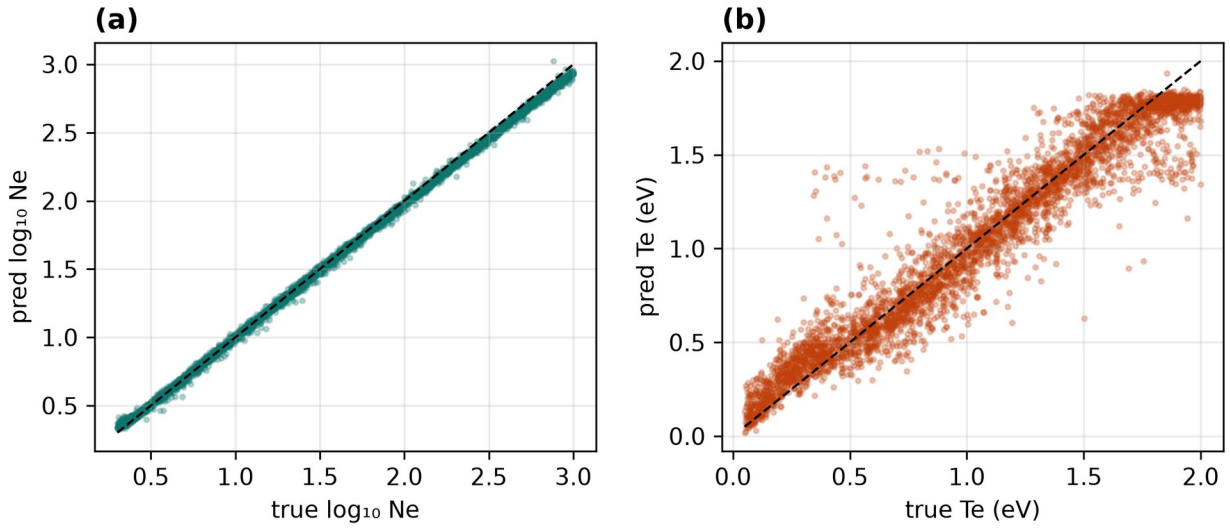


Figure 4. *Network prediction against the known true value on 4,000 held-out simulated sweeps for (a) density and (b) temperature. Density is recovered almost exactly; temperature is noisier and flattens near the top of the training range, around 1.8 eV.*

3.1.3 Calibration of the error bars

Every recovered value comes with an error bar from Monte-Carlo dropout (Gal & Ghahramani 2016): the network is run 50 times with random parts switched off, and the spread of its answers is the error bar. We checked those error bars against the known truth on the 4,000 simulated test sweeps. A well-calibrated error bar contains the true value about 68% of the time at one sigma and 95% at two sigma.

Density came out at 99.2% within one sigma, plasma potential at 96.8%, and temperature at 64.0% (89.4% within two sigma). So the density and plasma-potential error bars are about three times wider than they need to be, which is conservative. The temperature error bars are about the right size, but miss a little more often than they should in the tails.

On the real archive there is no truth, but the two halves of each commanded sweep should differ by no more than their combined error bars. They do: 99.9% of 9,895 pairs fall within one sigma. This shows the bars are not too narrow, but it cannot show whether they are too wide: the network's uncertainty about its mapping is nearly the same for both halves and cancels in their difference. The error bars are therefore calibrated on simulated sweeps only, and their calibration on real data is not known.

3.2 Validation on the real archive

3.2.1 Agreement with the classical method

Table 5 compares the two methods on the 15,147 real sweeps where both work and the instrument did not clip. The median difference shows any systematic offset; the typical difference is the median of the absolute difference, sweep by sweep.

Table 5. *Agreement between the network and the classical fit on the 15,147 unclipped sweeps both methods fit.*

Quantity	Median difference	Typical per-sweep difference
Electron density	−0.0061 dex (1.4%)	4.4%
Electron temperature	+0.037 eV	0.091 eV (25%)

There is no meaningful offset in density (1.4%), and individual sweeps differ by 4.4%, between a network trained only on simulated data and an independent curve-fitting procedure. This is the strongest evidence in this report that the network is not responding to an artefact of the preprocessing. Temperature agrees less well. The offset is small (+0.037 eV), but individual sweeps differ by 0.091 eV, or 25%, and by more on ADC channel 2 (0.111 eV) than on channel 1 (0.084 eV). In 2,331 of these sweeps (15.4%) the classical fit places Vp below Vf; without them the typical difference is 3.9% in density and 0.088 eV in temperature.

3.2.2 Consistency between the two halves of each sweep

Every commanded sweep is a triangle: the voltage goes up and then comes back down. The two halves are separate measurements of the same plasma a few seconds apart, so they should agree (Table 6).

Table 6. *Agreement between the up and down halves of each commanded sweep.*

Quantity	Value	Note
Pairs compared	9,895	both halves accepted by the network
Median density disagreement	0.0104 dex	2.4%
Systematic offset, down minus up	−0.0015 dex	0.4%, no directional bias
Correlation between halves	0.992	$\log_{10}$ density

This 2.4% is the precision to quote with densities from this archive. It comes from repeated measurement rather than from a model.

3.2.3 The recovered sweeps

The checks above use sweeps the classical fit already handles. The recovered sweeps need their own check, because the R^2 condition that admits them is easier to pass than the classical test (Section 2.5).

Each commanded sweep offers a direct test. For 2,157 of the 5,115 recovered sweeps, the other half of the same commanded sweep, recorded a few seconds apart, passes the classical fit and is not clipped. Against that classical value the recovered density differs by 5.3% (median absolute difference; offset −1.2%) and the temperature by 0.100 eV (offset +0.006 eV). Where both halves pass the classical fit, the same comparison gives 4.7% and 0.093 eV, so the recovered sweeps are as consistent as the sweeps the classical fit handles itself. The other 2,958 recovered sweeps have no classical twin. On simulated sweeps given their parameters and resistor settings, the network's error is 4% in density and 23% in temperature, close to Table 4.

The same test cannot show what clipping does, because the twin of a clipped sweep is almost always clipped too. So real sweeps were clipped on purpose: for 3,000 unclipped sweeps that pass the classical fit, the curve was cut flat at 90%, 70% and 50% of its peak and passed through the network again. Real clipped sweeps go flat from a median of +1.6 V. Cut at 70% of the peak, flat from +2.9 V, the network's temperature falls to about a third of its value on the original curve (median −0.24 eV) and its density rises by 10%; cut at 90%, flat only above +9.0 V, temperature falls by 20% and density does not change. The classical fit fares no better on the same curves: cut at 70%, half of them (51%) still pass its acceptance test, with density 34% too low; cut at 90%, 87% pass with density 6% too low. Mild clipping therefore costs the network little, and the 168 mildly clipped sweeps that do have an unclipped classical twin agree with it to 3% in density. Excluding every clipped sweep is the conservative choice; without the clipping condition, 3,230 more sweeps would have been counted as recovered.

Applying the classical test itself to the network's parameters is stricter. In the classical retardation window, the slope implied by the network's temperature fits ln(I) with R^2 of at least 0.90 on 73% of the unclipped sweeps the classical fit passes, but on only 22% of the recovered ones, although on both the network's temperature agrees with the local slope on average (ratio 1.04 and 1.03). The retardation region of a recovered sweep is often not a clean exponential, which is largely why the classical fit rejected it, so recovered temperatures should be used with more caution than recovered densities. Figure 5 shows a typical case.

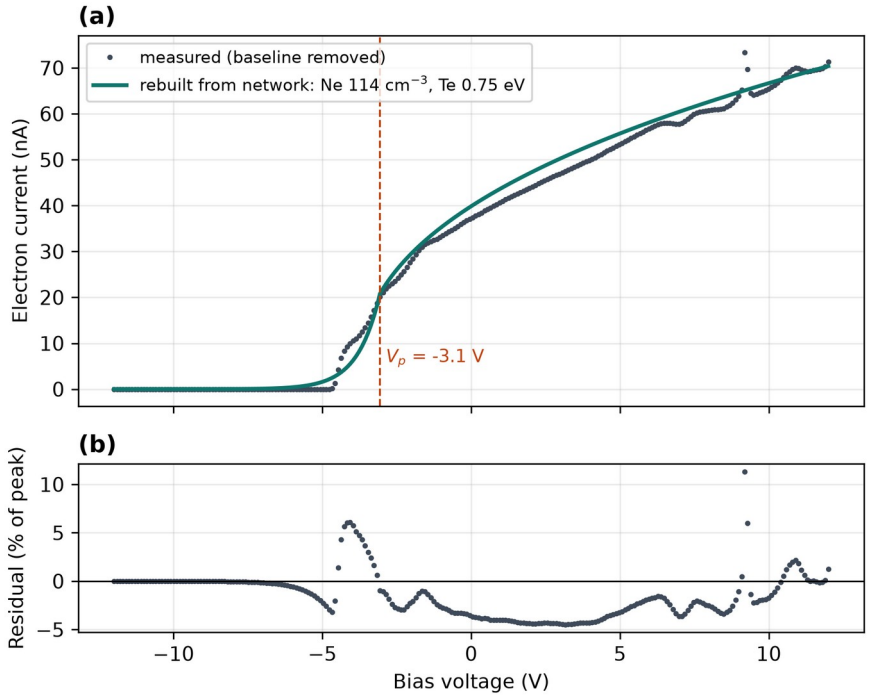


Figure 5. *A sweep rejected by the classical fit (retardation-fit R^2 0.79) and recovered by the network, from 25 August 2023 on ADC channel 2 at 1 MΩ. (a) Measured electron current and the curve rebuilt from the network's parameters (Ne 114 cm^{-3}, Te 0.75 eV, Vp −3.1 V; R^2 0.994). (b) Residuals: the largest misfit is just below the plasma potential, where Section 3.3.3 locates the forward model's limitation.*

3.2.4 Consistency with the mission timeline

The network sees only a current curve and is never told the date, the time or the solar geometry. Its daily medians nevertheless follow the mission timeline (Table 7). The classical fit shows the same

rising trend, so this checks that the network reads the day-to-day changes in the data consistently; it would not expose an instrumental effect shared by both methods.

Table 7. *Median recovered density per day.*

Date (2023)	Density (cm^{-3})	Date (2023)	Density (cm^{-3})
24 August	245	30 August	500
25 August	327	31 August	490
26 August	381	1 September	498
27 August	452	2 September	466
28 August	478	3 September *	99
29 August	488		

Density rises smoothly until about 29 August and then levels off near 490 cm^{-3}; the rank correlation between date and density is 0.78 (Spearman, n = 10 days, p = 0.008). The classical medians rise in the same way, from 102 to 526 cm^{-3} (ρ = 0.81, p = 0.005). (* 3 September is the terminator session; only 4 of its 408 sweeps were recovered, far too few to draw anything from, and it is excluded from the statistic.)

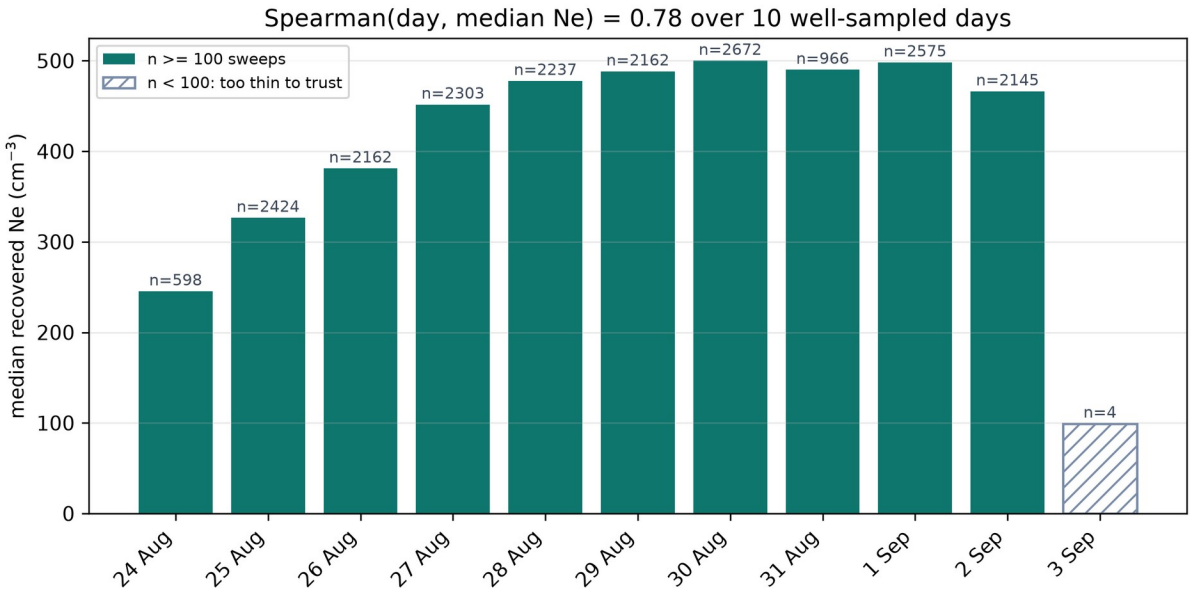


Figure 6. *Median recovered density per day. The network is never told the date. 3 September (hatched) has only 4 recovered sweeps and is left out of the trend.*

3.2.5 Comparison with the published analysis

The instrument team published their own analysis of this data (Manju et al. 2025), reporting mean daytime densities of 380–600 cm^{-3} and temperatures of 3,000–8,000 K. The median temperature of sweeps passing the network's acceptance test (0.38 eV, about 4,450 K) lies inside that range, although 16% of those sweeps fall outside it (12% below 3,000 K and 4% above 8,000 K). The daily median densities lie inside the published range from 27 August onwards (Table 7); the first three days are lower, as discussed in Section 3.4. For 2 September, Manju et al. (2026) give a mean density

of 478 cm^{-3} (range 409–680) for 19–20 h, before the hop. Over the same hour the network gives a mean of 433 cm^{-3} from 92 sweeps and the classical fit 423 cm^{-3} from 80 unclipped sweeps, both inside that range and 9–12% below its mean; on the three reference setups of Section 3.2.6 alone they are 447 and 430 cm^{-3}. The network's median plasma potential for the day is −3.8 V, against a published −3.9 V; the classical fit also gives −3.9 V. These are consistency checks rather than independent validation, because the sheath exponent was first chosen partly to match the published 2 September values, although Section 2.3 shows that the data alone give the same value.

3.2.6 Calibration between the probe setups

The instrument uses its six combinations of ADC channel and probe resistor in turn, one per sweep segment, so neighbouring segments of the same file measure the same plasma a few seconds apart through different setups. If every setup were calibrated alike, they would return the same density. Table 8 compares each sweep's density with the median classical density of the unclipped classical passes on the three low-resistance setups within six segments in the same file. Those three setups agree with each other to about 1%, with 94–98% of sweeps within 10%, and serve as the reference.

Table 8. *Density on each channel and resistor setup relative to the three low-resistance reference setups, measured on neighbouring segments of the same file (median ratio; in brackets, the share of sweeps within 10%). Clipped sweeps are excluded.*

Setup	Usable sweeps	Classical density vs reference	Network density vs reference	Note
Channel 1, 49.9 kΩ	4,839	1.00 (98%)	1.08	reference
Channel 2, 49.9 kΩ	4,855	0.99 (94%)	1.04	reference
Channel 1, 1 MΩ	4,876	1.01 (97%)	1.03	reference
Channel 1, 20 MΩ	4,433	0.83 (8%)	0.81	1.00 on 24 Aug, 0.81–0.89 from 25 Aug
Channel 2, 1 MΩ	1,259	0.46–1.00	0.44–1.00	0.54 on 24 Aug, 0.46 on 25–26 Aug, 1.00 on 1–2 Sep; 73% clipped
Channel 2, 20 MΩ	0	-	-	all but one sweep clipped; cannot be checked

Two setups read low. Channel 1 at 20 MΩ agrees with the reference on 24 August and reads 11–19% low from 25 August onwards, and its classical temperature is 1.45 times the reference value, against 0.91–1.15 among the reference setups themselves. Channel 2 at 1 MΩ reads about half the reference density from 24 to 26 August but agrees with it on 1–2 September. Because the setups sample the same plasma seconds apart, these offsets are instrumental; their cause, such as a gain or resistor-value error, was not established. Both methods show them. On channel 2 at 1 MΩ the network's recovered densities are at 0.54, 0.44 and 1.00 of the reference on the same days, and on channel 1 at 20 MΩ at 0.77–0.82 from 27 August, against 0.81–0.86 for the classical fit; only on 24–25 August does the network read high on that setup (1.30 and 1.11). The offset is therefore in the data rather than in the network. The checks of Sections 3.2.1–3.2.3 cannot see it, because each compares results from a single setup.

4,996 of the 20,262 usable sweeps (16.7% of the archive) lie on a setup and day whose density differs from the reference by more than 10%, and they are flagged in the output tables; their densities should be corrected or left out of absolute comparisons. They are flagged rather than corrected because the offset may not be a single factor: on channel 1 at 20 MΩ the ratio falls from 0.89 to 0.81 as the reference density rises, although days and density levels are mixed in that comparison. Without them, 51.1% of the archive is usable with no setup correction, and 48.8% comes from the three reference setups alone. The daily medians of Table 7 change little: on the reference setups alone they are 2–6% higher (257 to 506 cm^{-3}) and follow the same rise.

The comparison also gives the network's own error on real recovered sweeps with no offset mixed in. On the reference setups, recovered densities are 4.5% above their classical neighbours (median), with a typical difference of 6.1% and 58% within 10%, against 94–98% for classical passes. That is close to the 5.3% found against twin halves (Section 3.2.3), and it shows a small upward bias on recovered sweeps that the comparison on sweeps both methods fit (Section 3.2.1) does not.

3.3 Experiments

This section reports four further tests, including those that contradicted our initial expectations.

3.3.1 Effect of the physics term

The question. Physics-informed networks are usually evaluated on problems with a known solution (Raissi et al. 2019). Because we had an independent method and a self-consistency check, we could also measure the effect of the physics term on real data. We retrained the network with the physics requirement at four weights, from full strength down to switched off, three times each with different random starting points.

Table 9. *Effect of the physics-loss weight on simulated and real-data performance (mean of three training runs per weight).*

Physics weight	Simulated accuracy	Real self-consistency	Density offset vs classical	Real sweeps recovered
0.5 (what we use)	0.075 eV	0.0102 dex	−0.001 dex	39.4%
0.25	0.068 eV	0.0113 dex	+0.003 dex	39.5%
0.1	0.059 eV	0.0115 dex	+0.006 dex	39.1%
0.0 (switched off)	0.056 eV	0.0145 dex	+0.022 dex	38.8%

Result. Switching the physics term off improves the network on simulated data, with temperature error falling from 0.075 to 0.056 eV, but worsens it on the real archive in two ways. The two halves of each sweep disagree 42% more (0.0145 against 0.0102 dex), and a density offset of about 5% (+0.022 dex) appears against the independent classical fit, where with the term there is none. The number of recovered sweeps does not change measurably (39.4% against 38.8%).

Judged on simulated accuracy alone, the physics term would have been removed, leaving the network's real densities less repeatable and biased against an independent method. The term exists to help the network cope with the gap between simulation and reality, which a test run entirely in simulation cannot see.

We also measured how small a difference this comparison can resolve. Between the three runs of each weight, the recovered count varies by 40 to 120 sweeps (standard deviation), about one percentage point, so none of the recovery differences in Table 9 is measurable. The self-consistency differences are: the run-to-run spread is 0.0001 to 0.0008 dex, so 0.0102 dex at weight 0.5 is clearly better than 0.0113 dex at 0.25 and 0.0145 dex with the term off, and the weight was kept at 0.5. An earlier version of this comparison, made before clipped sweeps were excluded, found a ten-point gain in recovery from the physics term; that gain came from clipped sweeps and is withdrawn.

3.3.2 Testing for a second electron population

The question. A second, hotter electron population was physically plausible at this site, because sunlight releases photoelectrons from the surface (Poppe & Horányi 2010) and from the lander itself. We are not aware of a machine-learning Langmuir inversion that separates two populations, so we built one and tested the idea directly.

We extended the network from three outputs to five.

A five-output model will always report some second population, so the test is whether it does so when none exists. We trained it on data containing two populations and then applied it to curves built with only one.

The first version reported a hot population of 8.9% of the density on curves that had none. The cause was the training data: the hot population's size was drawn from a range that never included zero, so the network had never seen an example with no hot population.

Making 35% of the training examples single-population brought the false rate down to 1.7%, and only then was the model applied to real data.

The earlier evidence for two populations in this archive came from a common diagnostic: averaging many sweeps and examining how the slope varies across the curve. A single population gives a constant slope, and an early version of this analysis showed about a five-fold variation in local temperature. Its code was not saved and the records of its values disagree, so that figure cannot be reproduced; a rerunnable diagnostic gives local temperatures of 0.46–0.69 eV on the averaged archive.

A varying slope is not decisive. Each sweep has its own temperature, and an average of curves that fall off at different rates does not fall off at a single rate: it bends, even when every individual sweep has one population.

The effect is large here. Classical-fit temperatures run from 0.24 to 1.61 eV (1st to 99th percentile), and the deep part of an averaged curve is dominated by its hottest members: the hottest 10% of sweeps contribute 69% of the signal in the region where the inference is made. These and the tests below use unclipped sweeps only.

We built a control by giving each sweep a single population at its own measured temperature and averaging the results identically. It produces a graded slope with no second population present: local temperature falls from 0.78 eV at 1.9 V below the plasma potential to 0.42 eV at 0.3 V.

With that control in place we tested the archive three separate ways, and none supports a second population:

1. Against the matched control, the region where a hot population would dominate shows no consistent excess: at 1.9 and 1.5 V below the plasma potential it is positive on only 3 of the 8

days with enough sweeps to compare, at both depths. Closer to the plasma potential the real curves are flatter than the control, on 8 of 8 days at 0.3 V, which Section 3.3.3 attributes to the forward model rather than to a hot population.

2. A per-sweep test avoids averaging by measuring the slope in two voltage windows of the same sweep, which must agree for a single population. On simulated data it detects a hot population of 5% of the density every time.
3. On real sweeps that test gives 33.8% against a control's 34.7%, a difference of 0.8 sigma, which is not significant.

Splitting that result by probe channel showed the two channels disagreeing at 11 sigma, in opposite directions, so the small overall difference was two opposing effects cancelling.

This also argues against a second population. A real population in the plasma would affect both channels in the same way, so a signal that reverses sign between them points to the instrument.

Conclusion. This archive does not require a second electron population, which agrees with the published analysis of the same data. An earlier draft of this project claimed otherwise, with specific numbers; those are withdrawn and listed in Supplementary Appendix A.

3.3.3 Limits on temperature accuracy

The question. Density is recovered to about 5% and plasma potential to about 0.14 V, but temperature error stays near 20%. We examined whether this is a fundamental limit of the instrument or could be reduced.

Table 10 sets this against the MAVEN Langmuir probe.

Table 10. *Electron temperature accuracy reported for MAVEN and for this work.*

Mission / study	Temperature error	What kind of number
MAVEN at Mars (Ergun et al. 2015)	better than 20%	stated for most conditions
This work	**20–25%**	**simulation plus real repeatability**

An error near 20% is comparable with the figure stated for MAVEN, although that is an absolute accuracy, whereas ours comes from simulation and repeatability (Section 3.4).

A Fisher-information calculation for this measurement (450 cm^{-3}, 0.4 eV, quiet-channel noise) gives a Cramér–Rao bound of about 1.5% on temperature. The achieved error is about 13 times larger, so noise is not the limit.

The same calculation shows why. Temperature and plasma potential are 92% correlated: a slightly hotter plasma and a slight voltage shift bend the curve in nearly the same way. Not knowing the plasma potential inflates the achievable temperature error by a factor of about 2.5.

The training distribution also contributed. The network's absolute error was roughly constant across the temperature range (0.05–0.09 eV in every bin), and 41% of the training examples were hotter than 1.2 eV, a regime rarely seen in this archive, while only 10% were in the range it mostly occupies.

Training on the logarithm of temperature and narrowing the training range to realistic values gave 22.3% before and 20.2% after, over three runs each. The 1.5 sigma difference is not significant,

although results became more stable between runs. This comparison was run during development, and its script was not kept.

A more effective change removes the voltage from the problem before the network sees the curve. The curve is shifted so that zero sits at the plasma potential and divided by the current there so that density cancels, and its logarithm is taken. The slope of the result is set by the temperature alone.

Table 11. *Temperature error on simulated sweeps with and without the plasma-potential-aligned refinement.*

Configuration	Temperature error	Spread over 3 runs
As built (raw curve)	22.3%	± 1.3%
Aligned, using the true voltage (ceiling)	11.7%	± 2.6%
Aligned, using the predicted voltage	**12.2%**	**± 0.2%**

The error nearly halved, a difference of about 13 sigma over three independent runs in the 0.30–0.50 eV band, with every row at the full training budget (20,000 sweeps, 90 epochs). Using a predicted voltage instead of the true one costs little.

The change was then tested on real data. Following Section 3.3.1, it was not adopted on the strength of a simulated result; instead we measured whether the two halves of the same real sweep agreed better, a repeatability test that needs no ground truth. Table 12 gives the result, and the change fails it. The baseline in Table 12 is the final trained network rather than the three-run average of Table 11, so the simulated gain is smaller there: 0.075 to 0.049 eV, or 18.2% to 11.8% within the temperature range the archive mostly occupies.

Table 12. *The temperature refinement on simulated and real data. All values use a single pass with dropout switched off; clipped sweeps are excluded.*

	As built	With the fix
Simulated accuracy	0.075 eV	0.049 eV (better)
Real repeatability	**0.0097 eV**	**0.0164 eV (69% worse)**
Offset vs classical method	+0.009 eV	−0.040 eV (bias appeared)

The test itself is sound. A model predicting the same temperature for every sweep would score perfect repeatability while being useless, but this one's answers span 0.23 to 0.57 eV across the unclipped archive, and on the 6,405 unclipped sweeps whose two halves both pass the classical fit, its median difference between halves, with dropout switched off as in Table 12, is 0.006 eV against the classical fit's 0.025 eV. With the 50-pass average used elsewhere in this report the network's figure is 0.021 eV, because averaging a finite number of random passes leaves its own scatter.

We first suspected the region where the fix reads the temperature, near the plasma potential, where the real curves depart from the physics model on 8 of 8 days. Moving the reading window deeper made every measure worse, including the simulated accuracy, which rules this explanation out.

The current falls exponentially below the plasma potential. Just below it there is plenty of signal, but that is where the physics model is wrong; deeper down the model holds, but the current has dropped to 0.15% of its peak, below the noise floor on every channel.

The temperature information and the model's error therefore occupy the same part of the curve, and no window has one without the other.

Conclusion. Temperature accuracy here is limited by how well our physics model describes the curve near the plasma potential - not by the network. Further changes to the network are unlikely to improve it. Fixing that part of the physics model is the prerequisite, and it is the clearest piece of future work.

3.3.4 Per-sweep scatter

The per-sweep test in Section 3.3.2 produced something unrelated to electron populations. The temperatures it measures, from straight-line fits in two windows of a single sweep, scatter about 3.7 times more on real sweeps than our noise model predicts. The scatter is symmetric, so it is not a bias, but we cannot account for it. The same excess scatter appears independently when the up and down halves of each sweep are compared.

The practical consequence is that a temperature read from a local fit to a single sweep is much less repeatable than a simulated accuracy figure suggests, and should always be quoted with that scatter. Whether it comes from the baseline subtraction, from the plasma genuinely changing during a sweep, or from something the physics model omits, we could not determine.

3.4 Discussion

The recovered values broadly agree with the published analysis of the same instrument (Section 3.2.5). The exception is density on the first three days (245–381 cm^{-3}), which is below the published 380–600 cm^{-3}. On those days the network's median plasma potential is also far more negative (−8.2 to −4.7 V). Both analyses start from the same photoemission correction, which the archive applies with the method of Johansson et al. (2017), so this step is unlikely to explain the difference on its own, and the cause was not established here. The setup offsets of Section 3.2.6 explain only a small part of it: on the reference setups alone the first three days are still 257–401 cm^{-3}. Manju et al. (2026) also report densities of about 1,100 cm^{-3} near the terminator after the hop, but this could not be tested here because only 4 of the 408 sweeps from 3 September were recovered.

The accuracy figures should be read with their source in mind. The temperature error of about 20% on simulated data is similar to the better-than-20% stated for MAVEN in most conditions (Ergun et al. 2015), but that is an absolute accuracy, whereas the figures in this report come from simulation and from repeatability. The 2.4% disagreement between sweep halves is the most direct error estimate available for density, but it measures precision, not absolute accuracy.

The absence of a systematic offset against the classical fit shows that the network is not simply reproducing artefacts of preprocessing. It does not remove errors that the two methods share, however. Both rely on the OML picture with a fixed sheath-expansion exponent. Liu et al. (2023) found that assuming the standard OML exponent for their needle probes overestimated density by a factor of about three. A shared error of this kind would move both methods together and would not appear in any of the consistency checks used here.

The physics-term experiment (Section 3.3.1) has a practical lesson for physics-informed machine learning. Physics-informed networks are usually judged by their accuracy on problems where the

exact solution is known (Raissi et al. 2019). Here, accuracy on simulated sweeps pointed one way and the consistency of real sweeps the other. When the physics term is meant to help the model cope with the gap between simulation and reality, it needs to be judged on real data.

For the two-population test, the physical expectation came from photoemission. A sunlit surface produces a photoelectron sheath (Poppe & Horányi 2010), and the lander and probe add photoelectrons of their own, so a warmer secondary population near the probe was plausible. The controlled tests in Section 3.3.2 do not support it, which is consistent with the published single-population analysis (Manju et al. 2025). The wider point is about method: when sweeps with different temperatures are averaged, the averaged curve shows a changing slope like that of a two-population plasma. Any claim of multiple populations based on averaged sweeps therefore needs a matched single-population control.

Finally, the temperature experiments show where further progress has to come from. The part of the curve that carries most of the temperature information, just below the plasma potential, is also where the forward model departs from the real sweeps. Better network design cannot get around this; a better description of the current near the plasma potential, including a variable sheath exponent and an explicit photoemission term, is needed first.

4. Limitations and future scope

4.1 Limitations

- We cannot check our absolute accuracy. There is no second instrument on the lander and no lunar equivalent of ground radar. Everything in this report is consistency, not verified accuracy.
- The physics model assumes a fixed value for how the collecting region grows with voltage. A study of cylindrical needle probes shows that this matters: assuming the standard value overestimated density by a factor of about three (Liu et al. 2023). This is the largest known systematic remaining.
- Photoemission from the probe and lander is not part of the forward model. The rawB product already has a photoemission current removed, estimated with the method of Johansson et al. 2017 as in the published analysis (Manju et al. 2025), and any remaining constant is absorbed by the baseline subtraction. Photoemission that changes with time is not corrected.
- A channel-dependent effect distorts the curve near the plasma potential. The two probe channels disagree at 11 sigma in the per-sweep test (Section 3.3.2). It persists when clipped sweeps are removed. We do not know its cause, and it limits how far any shape-based conclusion can be pushed.
- 30% of the sweeps are clipped by the instrument and are excluded from both methods, because the forward model has no saturation and neither method reads them correctly. Most are on ADC channel 2, the same channel that shows the effect near the plasma potential described above, so the two may share a cause in that channel's front end. Mildly clipped sweeps could probably be kept: the 442 on channel 1 at 20 MΩ, flat only above about +11 V, match the unclipped sweeps on that setup, and admitting them would give 69.3%. This was not tested in general.
- Two of the six channel and resistor setups read density low against the setups measuring the same plasma: channel 1 at 20 MΩ by about 17% from 25 August, and channel 2 at 1 MΩ by about half on 24–26 August (Section 3.2.6). The cause is not known, and the 4,996 usable

sweeps affected are flagged rather than corrected. Temperature also differs between setups, by up to about 25% even among the three reference setups (0.91 to 1.15 of their common value), so absolute temperatures carry a setup-dependent uncertainty of that size.

- The error bars are calibrated on simulated sweeps only. On real data the comparison of sweep halves shows they are not too narrow, but cannot show whether they are too wide.
- Recovered temperatures are less secure than recovered densities: the retardation region of many recovered sweeps does not follow a single exponential (Section 3.2.3).
- The 3 September terminator session is still almost entirely unrecovered - 4 sweeps of 408.
- The hop itself could not be tested. The project plan included a blind test on the lander's 2 September hop, where the instrument team report density falling from about 478 to 26 cm^{-3} (Manju et al. 2026). During the hop the probe was in a fixed-voltage mode and recorded no sweeps at all, so there is nothing for any sweep-based method to analyse.
- We tried pre-training on simulated Mars-like data before fine-tuning on lunar data. It made the model about fifteen times worse and was abandoned. We report this as a negative result about this particular attempt, not as a general claim: both datasets came from the same physics model, so the pre-training carried no new physics, only a prior centred two to three orders of magnitude away in density.

4.2 Future scope

Table 13. *Future work, with the reason for each item.*

Task	Why
Fix the physics model near the plasma potential	Currently limits temperature accuracy (Section 3.3.3)
Let the sheath-growth exponent vary per sweep	Shown to matter for cylindrical probes (Liu et al. 2023); largest known systematic
Add photoemission as a fitted term	It is present physically and absent from the model
Identify the channel-dependent effect	Unexplained, and it bounds shape-based conclusions
Calibrate the six setups against each other	Two read density low by 17% to 54% (Section 3.2.6)
Simulate instrument clipping in training	30% of sweeps are clipped and currently unusable

5. Conclusions

This study shows that a neural network trained only on simulated sweeps can invert real RAMBHA-LP Langmuir probe measurements from the lunar surface. On the 15,147 unclipped sweeps that the classical fit can also handle, the two methods show no meaningful density offset and differ by 4.4% per sweep. The network recovers 5,115 of the 5,641 unclipped sweeps that the classical fit rejects (91%), raising the usable part of the archive from 51% to 68%; where a recovered sweep has a classical twin half, the two densities agree to 5.3%. The rising and falling halves of each sweep agree to 2.4% in density, the daily medians follow the same rise across the mission as the classical fit, and the densities for 2 September lie inside the range published by the instrument team. About 30% of the archive is clipped by the instrument and cannot be read by either method. Two of the six channel

and resistor setups also read density low against the others, by about 17% and, early in the mission, by about half; both methods share the offset, and 51% of the archive is usable without any setup correction.

The physics-informed loss term lowers accuracy on simulated data, with a temperature error of 0.075 eV against 0.056 eV without it (three runs each), but on real sweeps it reduces the disagreement between sweep halves by 30% and removes a 5% density offset against the classical fit. It does not change how many sweeps are recovered. Judged on simulated accuracy alone, it would have been removed. Choices that are meant to close the gap between simulation and reality must therefore be tested on real data.

A two-population version of the network, checked against a single-population control, finds no evidence that the archive requires a second electron population. The remaining signal reverses sign between the two probe channels, which points to an instrumental cause. The local slope of an averaged sweep is not, on its own, evidence for multiple populations, because a spread of single-population temperatures also gives a graded pattern; in this archive the hottest 10% of sweeps supply 69% of the signal in the region used for that inference.

Electron temperature remains the least certain quantity, at about 20% on simulated data and 25% per sweep against the classical fit. An input aligned to the plasma potential reduces this error by between a third and a half in simulation but makes the real repeatability 69% worse, because the region of the curve that carries the temperature information is also where the forward model departs from the data. Improving the physics model near the plasma potential is therefore the next step, before any further change to the network.

Acknowledgements

The RAMBHA-LP data used in this work were provided by the Indian Space Research Organisation through the ISRO Science Data Archive (PRADAN), operated by the Indian Space Science Data Centre (ISSDC). The RAMBHA-LP instrument was built and operated by the Space Physics Laboratory, Vikram Sarabhai Space Centre. The author thanks Dr. Rama Chandra Pillutla, Lab for Spatial Informatics, IIIT Hyderabad, for supervising this independent study and for his guidance throughout the work.

Data and code availability

The RAMBHA-LP data are available from the ISRO Science Data Archive (ISRO PRADAN), release LTA_RAMBHA-LP_August2024_final. The analysis code, configuration and output tables are available at https://github.com/Suniljadaun/Chandrayaan3_Rambha-LP-ML.

References

Ergun, R. E., Morooka, M. W., Andersson, L. A., et al. (2015). Dayside electron temperature and density profiles at Mars: First results from the MAVEN Langmuir probe and waves instrument. *Geophysical Research Letters*, 42(21), 8846–8853. https://doi.org/10.1002/2015GL065280

Gal, Y., & Ghahramani, Z. (2016). Dropout as a Bayesian approximation: Representing model uncertainty in deep learning. *Proceedings of the 33rd International Conference on Machine Learning*, PMLR 48, 1050–1059. https://proceedings.mlr.press/v48/gal16.html

Hutchinson, I. H. (2002). *Principles of Plasma Diagnostics* (2nd ed.). Cambridge University Press. https://doi.org/10.1017/CBO9780511613630

ISRO Science Data Archive (PRADAN). Chandrayaan-3 RAMBHA-LP PDS4 data archive, release LTA_RAMBHA-LP_August2024_final. Indian Space Science Data Centre. https://pradan.issdc.gov.in/ch3/ (accessed 2026)

Johansson, F. L., Odelstad, E., Paulsson, J. J. P., et al. (2017). Rosetta photoelectron emission and solar ultraviolet flux at comet 67P. *Monthly Notices of the Royal Astronomical Society*, 469(Suppl. 2), S626–S635. https://doi.org/10.1093/mnras/stx2369

Liu, G., Marholm, S., Eklund, A. J., Clausen, L., & Marchand, R. (2023). m-NLP inference models using simulation and regression techniques. *Journal of Geophysical Research: Space Physics*, 128(2), e2022JA030835. https://doi.org/10.1029/2022JA030835

Manju, G., Pant, T. K., Mridula, N., et al. (2025). In situ ionospheric observations near lunar south pole by the Langmuir Probe on Chandrayaan-3 lander. *Monthly Notices of the Royal Astronomical Society*, 542(4), 2647–2656. https://doi.org/10.1093/mnras/staf1276

Manju, G., Pant, T. K., Mridula, N., et al. (2026). Langmuir probe-based lunar near-surface plasma measurements associated with the unique hop test of Chandrayaan 3 lander: observations and model simulations. *Monthly Notices of the Royal Astronomical Society*, 546(4), stag077. https://doi.org/10.1093/mnras/stag077

Mott-Smith, H. M., & Langmuir, I. (1926). The theory of collectors in gaseous discharges. *Physical Review*, 28(4), 727–763. https://doi.org/10.1103/PhysRev.28.727

Poppe, A., & Horányi, M. (2010). Simulations of the photoelectron sheath and dust levitation on the lunar surface. *Journal of Geophysical Research: Space Physics*, 115, A08106. https://doi.org/10.1029/2010JA015286

Raissi, M., Perdikaris, P., & Karniadakis, G. E. (2019). Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378, 686–707. https://doi.org/10.1016/j.jcp.2018.10.045